

Async vs. Sync Programming

Dmitry (Mitya) Komanov



dmitryk@wix.com



@dkomanov



linkedin/dkomanov



github.com/dkomanov

01

Terms

Synchronous (or blocking)

- Once blocking method called, calling thread will wait.
- Price: context switch

Asynchronous (or non-blocking)

- Thread continues to work after call to async method. Result of async method execution may be acquired via callback or via polling.
- Price: queue, poll, await.

02

That's it?

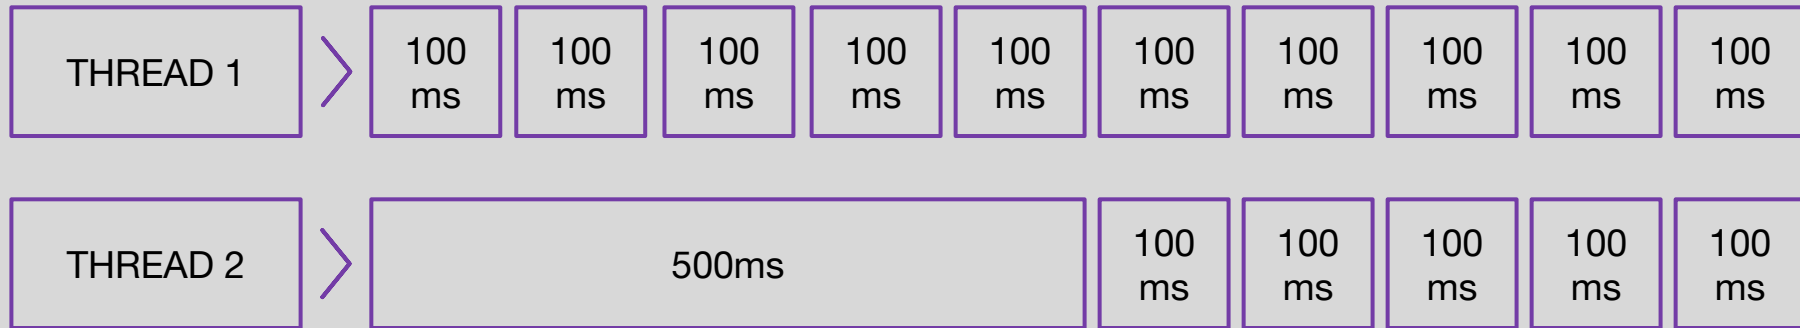
We're talking about I/O:
disk and network

Averaged flow



This is how 20 RPS looks like

Something happened flow

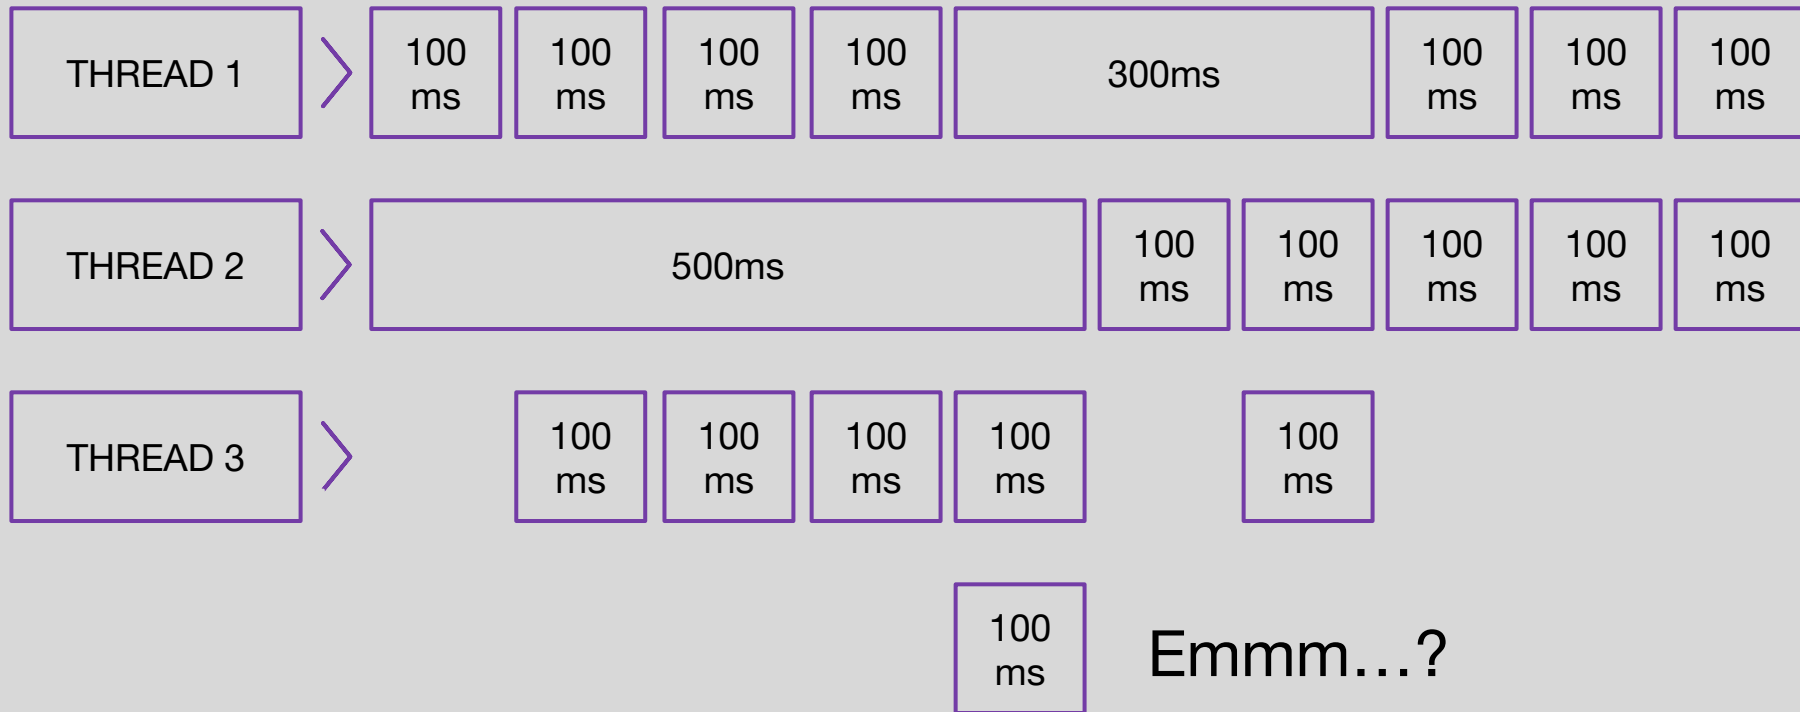


This is how 16 RPS looks like
Wait, where are 4 more RPS?

Add another thread



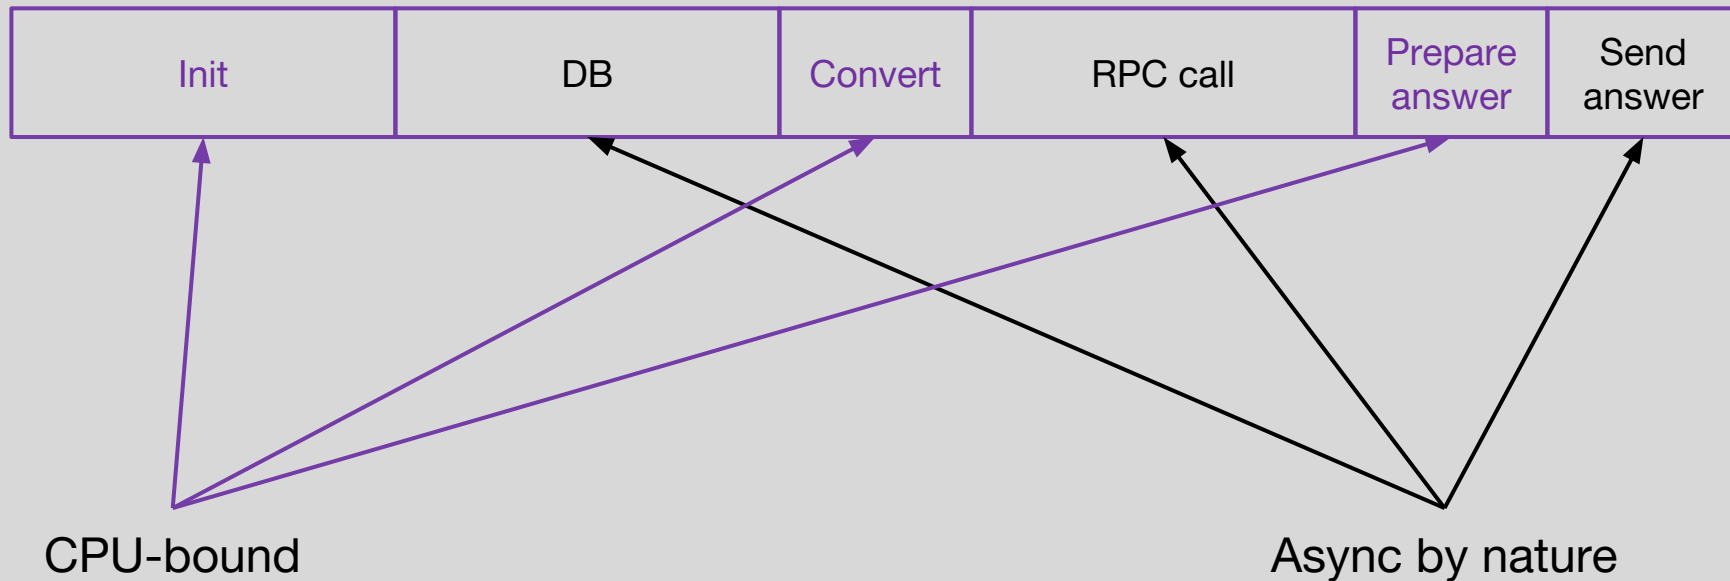
But... the real life is far from picture



03

Going deeper

Inside request



Context Switch is somewhat
expensive operation

Context Switch

- Default time slice is 100 msec
- Context Switch cost 3-30 usec
- 3+ orders of magnitude difference...
- ... but how much time do you spent for computations in your code?
- 3 ns access to L1 cache, 70 ns access to memory ¹⁶.
- Not to mention wait time till reschedule.

Thread is an expensive
thing by itself

Thread

- Thread allocates **X** KB for its stack (via -Xss option).
- By default on amd64 it's 1024 KB ⁴.
- Minimal value is 228 KB ⁵.
- 200 Jetty threads will consume 200MB non-heap memory ⁶ :)

Scaling in blocking system

- Adding more threads.
- More strict timeouts.
- Throttlers, circuit breakers etc.

But what if we still want to serve
some of that rejected requests?

04

Let's go async!

Async request processing



- Init and start async request to DB;
- Process another request;
- Once DB responded, convert data and issue an RPC call;
- Process another request;
- Once RPC responded, prepare answer and start sending.

Cons of async programming

- Request can be executed in different threads.
- Therefore, monitoring is more complicated.
- Harder to debug.
- Harder to control concurrency, resource usage. Therefore, harder to implement backpressure.
- Code is more complex.

Monitoring

- NewRelic supports limited amount of async libraries/frameworks *
- You can't just wrap piece of code into `try { meter } finally { report }` block, you have to pass start time along with request.
- Monitoring of queues for shared resources isn't optional anymore.

Debug

- Thread dump is almost useless. You will see polling threads without any context.
- Need to search by transaction id.
- Logs are essential (or any other trace tool).

Concurrency Control

- No more single Jetty thread pool to control concurrency on application.
- Need to configure all queues and make it granular (i.e. queue for each rpc client).
- Have to think about resource limits (for each resource and global limit).

Scaling in non-blocking system

- Fine-tuning of queue size and behavior (LIFO vs FIFO).
- More strict timeouts.
- Throttlers, circuit breakers etc.

All async-by-nature operations
already implemented async in kernel

05

Where is the gain?

No magic.
Sorry.

According to Netflix⁸

- You will get resilience for your system.
- You will get ability to handle more concurrent connections.
- You won't get performance boost (or maybe yes).
- You won't get latency improvement (or maybe yes).
- You will get more complicated code.

Async is like TDD
:)

06

Under the Hood

Based on:

- epoll (linux) ⁹ or
- kqueue (bsd, mac os) ¹⁰ or
- IOCP (windows) ¹¹

epoll¹²

- Single file descriptor aggregates many file descriptors.
- Return fds that are available to operate.
- Works for constant time (as opposite to $O(N)$).
- 1 syscall vs N syscalls.

07

Why to use sync?

You already use it.
Proven to work.

But really, why?

- Linear execution flow: easier code.
- Everyone got used to it.
- Sometimes you don't care.

08

When to use sync?

Use blocking if

- You don't care about handling a lot of requests.
- Your application is mostly CPU-bound.
- Your application read/write a lot from few files/sockets.
- You aren't afraid of OOM storms.

09

Async Programming

Use node.js, it's async
from the first day.

Event Loop¹³

This is how event loop can look like:

- You can't spawn threads
- You may only enqueue tasks to the queue
- Different tasks will be executed somehow in background and will issue an event once it's done.

```
function main() {  
  while (true) {  
    val task = taskQueue.take();  
    if (task.type === 'Function') {  
      task.execute();  
    } else if (task.type === 'BytesReadFromFile') {  
      task.callback.call(task.bytes);  
    } else if (task.type === 'BytesReadFromSocket') {  
      task.callback.call(task.bytes);  
    }  
  }  
}
```

Future/Promise

- Typically it's an abstraction over thread pool (but not necessarily).
- Callbacks!
- But you may use for comprehensions in Scala ¹⁴.
- Very important to restrain temptation of using Await.

Future/Promise

Future itself doesn't mean using async IO, but it allows you to think so.

In this example:

- We use `DeferredResult` which allow to return worker thread to jetty back (a kind of continuation).
- Decoupling of worker threads of jetty and application logic threads.

```
@RequestMapping("/action")
```

```
def action(): DeferredResult[Response] = {  
    val result = new DeferredResult[Response]()  
    val future: Future[Response] = executeActionAsync()  
    future.onComplete {  
        case Success(a) => result.setResult(a)  
        case Failure(e) => result.setErrorResult(e)  
    }  
    result  
}
```

Actor Model¹⁵

- Actor is an execution unit.
- Actor has an input queue (mailbox).
- Actor can send [immutable] messages to other actors.
- One actor can be executed only in one thread at any point of time (each actor is single-threaded).
- Therefore, actor can store mutable state, access to this state doesn't require synchronization.

Actor Example

This is simple example of actors:

- SumActor have a cache, it's safe to access it.
- Actor framework take care of execution in a single thread for each actor.

```
class Actor[T] {  
  private val queue: NonBlockingQueue[T] = ...  
  protected def execute(message: T): Unit  
  final def send(message: T) = queue.put(message)  
  final def run = // magic to get messages from queue and pass it to execute  
}  
  
case class SumMessage(a: Int, b: Int, nextActor: Actor[Int])  
  
class SumActor extends Actor[Sum] {  
  val cache: Cache[(Int, Int), Int] = ...  
  override def process(m: Sum): Unit = {  
    val sum = cache.getOrElseUpdate(m.a -> m.b, {  
      m.a + m.b  
    })  
    m.nextActor.send(sum)  
  }  
}
```

Thank You!



<https://goo.gl/t4Ack2>



dmitryk@wix.com



@dkomanov



linkedin/dkomanov



github.com/dkomanov

References

1. <http://stackoverflow.com/questions/16401294/how-to-know-linux-scheduler-time-slice>
2. <http://www.linuxplumbersconf.org/2013/ocw//system/presentations/1653/original/LPC%20-%20User%20Threading.pdf>
3. <http://blog.tsunanet.net/2010/11/how-long-does-it-take-to-make-context.html>
4. <http://www.oracle.com/technetwork/java/javase/tech/vmoptions-jsp-140102.html>
5. <http://xmlandmore.blogspot.com/2014/09/jdk-8-thread-stack-size-tuning.html>
6. <https://github.com/eclipse/jetty.project/blob/jetty-9.4.x/jetty-util/src/main/java/org/eclipse/jetty/util/thread/QueuedThreadPool.java>
7. <https://docs.newrelic.com/docs/agents/java-agent/getting-started/compatibility-requirements-java-agent>
8. <http://techblog.netflix.com/2016/09/zuul-2-netflix-journey-to-asynchronous.html>
9. <https://en.wikipedia.org/wiki/Epoll>
10. <https://en.wikipedia.org/wiki/Kqueue>
11. https://en.wikipedia.org/wiki/Asynchronous_I/O
12. <https://www.quora.com/Network-Programming-How-is-epoll-implemented>
13. <https://www.youtube.com/watch?v=8aGhZQkoFbQ>
14. <http://danielwestheide.com/blog/2013/01/09/the-neophytes-guide-to-scala-part-8-welcome-to-the-future.html>
15. https://en.wikipedia.org/wiki/Actor_model
16. <https://news.ycombinator.com/item?id=13934288>